

- [RPi Cluster - Swarm](#)
 - [Materials](#)
 - [Research and reference](#)
 - [Install Docker](#)
 - [Option 1 - Install Docker as Addon to Jessie](#)
 - [Option 2 - Install Hypriot \(used here\)](#)
 - [Reference](#)
 - [Use the utility program `flash` to download image and copy to SD](#)
 - [Create SSH Keys](#)
 - [Copy SSH Key of Master Node \(rpi-m1\) to Workers \(rpi-s1, rpi-s2, rpi-s3\)](#)
 - [Test SSH Key](#)
 - [Create Docker Swarm Manager](#)
 - [Reference](#)
 - [Architecture](#)
 - [Features](#)
 - [UML](#)
 - [Initialize the Swarm Manager at Master Node](#)
 - [Add 3 Worker Nodes to Join Swarm](#)
 - [Review Swarm Info](#)
 - [List of Nodes](#)
 - [Running Services in the Docker Swarm](#)
 - [Monitoring](#)
 - [Physical IO](#)
 - [Summary](#)
 - [Cheat Sheet](#)
 - [CLI](#)

RPi Cluster - Swarm

Materials

- 1 [USB power hub](#) with 4 ports 2amp per port minimun.
- 4 Raspberry Pi 2B (1GB, no wifi)
- 4 SD cards minimun 8GB, I use 32GB here.
- 4 wifi dongles (for rpi2)
- 1 router (to be used later)

Research and reference

- [Raspberry Pi 3 Super Computing Cluster Part 2 - Software Config](#) <<--- follow this for networking 20170522
- [Setup Kubernetes on a Raspberry Pi Cluster easily the official way!](#) <<--- follow this 20170522
- [How I setup a Raspberry Pi 3 Cluster Using The New Docker Swarm Mode In 29 Minutes](#)
- [Cluster computing on the Raspberry Pi with Kubernetes](#)
- [Docker Swarmmode Test Scenarios](#)
- [Get Started with Docker on Raspberry Pi](#) has additional notes about create docker image and access GPIO.

Install Docker

Option 1 - Install Docker as Addon to Jessie

- create a SD with Jessie (2017-04-10-raspbian-jessie.img) on it.
- [install docker](#)

```
$ curl -sSL https://get.docker.com | sh
```

- add current user, pi, to docker group to avoid type `sudo` each time.

```
$ sudo usermod -aG docker pi
```

- auto start

```
$ sudo systemctl enable docker
```

```
pi@rpi-01:~$ sudo systemctl enable docker
Synchronizing state for docker.service with SysVinit using update-rc.d...
Executing /usr/sbin/update-rc.d docker defaults
Executing /usr/sbin/update-rc.d docker enable
pi@rpi-01:~$
```

- this gives use rpi with
 - image 2017-04-10-raspbian-jessie.img
 - dynamic wifi (consider to move to static IP later)
 - docker

This requires installation of docker-machine and docker-compose separately. See procedures at the links below:

- [docker end to end app](#), this covers building OS images for ARM and GPIO applications.

Option 2 - Install Hypriot (used here)

Hypriot provides a distribution, Hypriot, which has docker pre-installed.

Reference

- [Docker on ARM Raspberry Pi](#)
- [\[Guide\] Installing Docker Swarm on HypriotOS](#)

hostname.local does not work out of box.

```
HypriotOS/armv7: pirate@rpi-m1 in ~
$ ping rpi-s2.local -c5
ping: unknown host rpi-s2.local
HypriotOS/armv7: pirate@rpi-m1 in ~
$ hostnamectl
  Static hostname: rpi-m1
        Icon name: computer
        Chassis: n/a
        Machine ID: 9989a26f06984d6dbadc01770f018e3b
        Boot ID: f1c4d9d1856343c3a15fcd68b3ad59bd
        Operating System: Raspbian GNU/Linux 8 (jessie)
        Kernel: Linux 4.4.50-hypriotos-v7+
        Architecture: arm
HypriotOS/armv7: pirate@rpi-m1 in ~
```

We need to install [mDNS service](#) for example, `avahi-daemon`. There are other ways.

Use the utility program `flash` to download image and copy to SD

There is an utility [flash](#) to simple the process of copying image and setup configuration.

- get image from [here](#)
- default user "pirate" (password "hypriot").
- create a config file: device-init.yaml as below;

```
hostname: rpi-m1    <!-- modify this for each node
wifi:
  interfaces:
    wlan0:
      ssid: "<ACCESS-POINT-ID>"
      password: "<PASSWORD>"
```

- download and create a SD with device config file; modify hostname for each node.

```
flash -c device-init.yaml https://github.com/hypriot/image-builder-rpi/releases/download/v1.4.0/hypriotos-rpi-v1.4.0.img.zip
```

- or copy it by following instruction [here](#) to clone image, then change node name or just re-download with different node name in.

```
→ rpi2cluster git:(master) X nmap -sP 192.168.1.0/24 | grep rpi-
Nmap scan report for rpi-s3 (192.168.1.81)
Nmap scan report for rpi-s2 (192.168.1.84)
Nmap scan report for rpi-s1 (192.168.1.86)
Nmap scan report for rpi-m1 (192.168.1.88)
```

- check docker installation.

```
→ rpi2cluster git:(master) X ssh pirate@192.168.1.88
The authenticity of host '192.168.1.88 (192.168.1.88)' can't be established.
ECDSA key fingerprint is SHA256:CMQc15FF7V/t7N4hfcTmS+/xIuQ00xFLWYezymsUE18.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '192.168.1.88' (ECDSA) to the list of known hosts.
pirate@192.168.1.88's password:
```

```
HypriotOS (Debian GNU/Linux 8)
```

```
The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.
```

```
Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.
```

```
HypriotOS/armv7: pirate@rpi-m1 in ~
```

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
HypriotOS/armv7:	pirate@rpi-m1	in ~				

```
$ docker -v
```

```
Docker version 17.03.0-ce, build 60ccb22
```

```
HypriotOS/armv7: pirate@rpi-m1 in ~
```

check `docker-compose` and `docker-machine`

```
HypriotOS/armv7: pirate@rpi-m1 in ~
```

```
$ docker-compose -v
```

```
docker-compose version 1.11.2, build dfed245
```

```
HypriotOS/armv7: pirate@rpi-m1 in ~
```

```
$ docker-machine -v
```

```
docker-machine version 0.9.0, build 15fd4c7
```

```
HypriotOS/armv7: pirate@rpi-m1 in ~
```

swarm needs to be installed separately,

```
$ swarm <--- this is wrong command, use `docker swarm init`
```

```
-bash: swarm: command not found
```

```
HypriotOS/armv7: pirate@rpi-m1 in ~
```

There is bug for ping,

```
$ ping rpi-s1 ping: icmp open socket: Operation not permitted HypriotOS/armv7: pirate@rpi-m1 in ~ $ ping 192.168.1.86 ping: icmp open socket
```

Create SSH Keys

There is some error about cloning, see [here](#) for solution.

- delete the ssh entry in `~/.ssh/known_hosts` before `ssh` login the node.

```
ssh-keyscan -t rsa github.com >> ~/.ssh/known_hosts
```

```

HypriotOS/armv7: pirate@rpi-m1 in ~
$ ssh-keygen -t rsa -C "pirate@rpi-m1"
Generating public/private rsa key pair.
Enter file in which to save the key (/home/pirate/.ssh/id_rsa):
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/pirate/.ssh/id_rsa.
Your public key has been saved in /home/pirate/.ssh/id_rsa.pub.
The key fingerprint is:
c7:61:7a:96:68:2e:10:4e:08:3a:7f:f6:a6:63:4a:37 pirate@rpi-m1
The key's randomart image is:
+---[RSA 2048]-----+
|.                    |
|.. .                |
|o . o      o       |
| o o .   = o        |
| . =   S *          |
|  o o o +           |
| . E + .            |
| . .o+ .            |
| .O..               |
+-----+

```

Copy SSH Key of Master Node (rpi-m1) to Workers (rpi-s1, rpi-s2, rpi-s3)

```

HypriotOS/armv7: pirate@rpi-m1 in ~
$ ssh-copy-id pirate@192.168.1.86
The authenticity of host '192.168.1.86 (192.168.1.86)' can't be established.
ECDSA key fingerprint is 3a:da:ef:3e:57:dd:dc:6a:44:8c:21:87:36:a6:5c:2b.
Are you sure you want to continue connecting (yes/no)? yes
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter out any that are already installed
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you are prompted now it is to install the new keys
pirate@192.168.1.86's password:

Number of key(s) added: 1

Now try logging into the machine, with:  "ssh 'pirate@192.168.1.86'"
and check to make sure that only the key(s) you wanted were added.
$

```

repeat it for rpi-s2, rpi-s3

Test SSH Key

```

HypriotOS/armv7: pirate@rpi-m1 in ~
$ ssh pirate@192.168.1.86

HypriotOS (Debian GNU/Linux 8)

The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.
Last login: Tue May 16 20:36:41 2017 from waterlily
HypriotOS/armv7: pirate@rpi-s1 in ~

```

Works.

Create Docker Swarm Manager

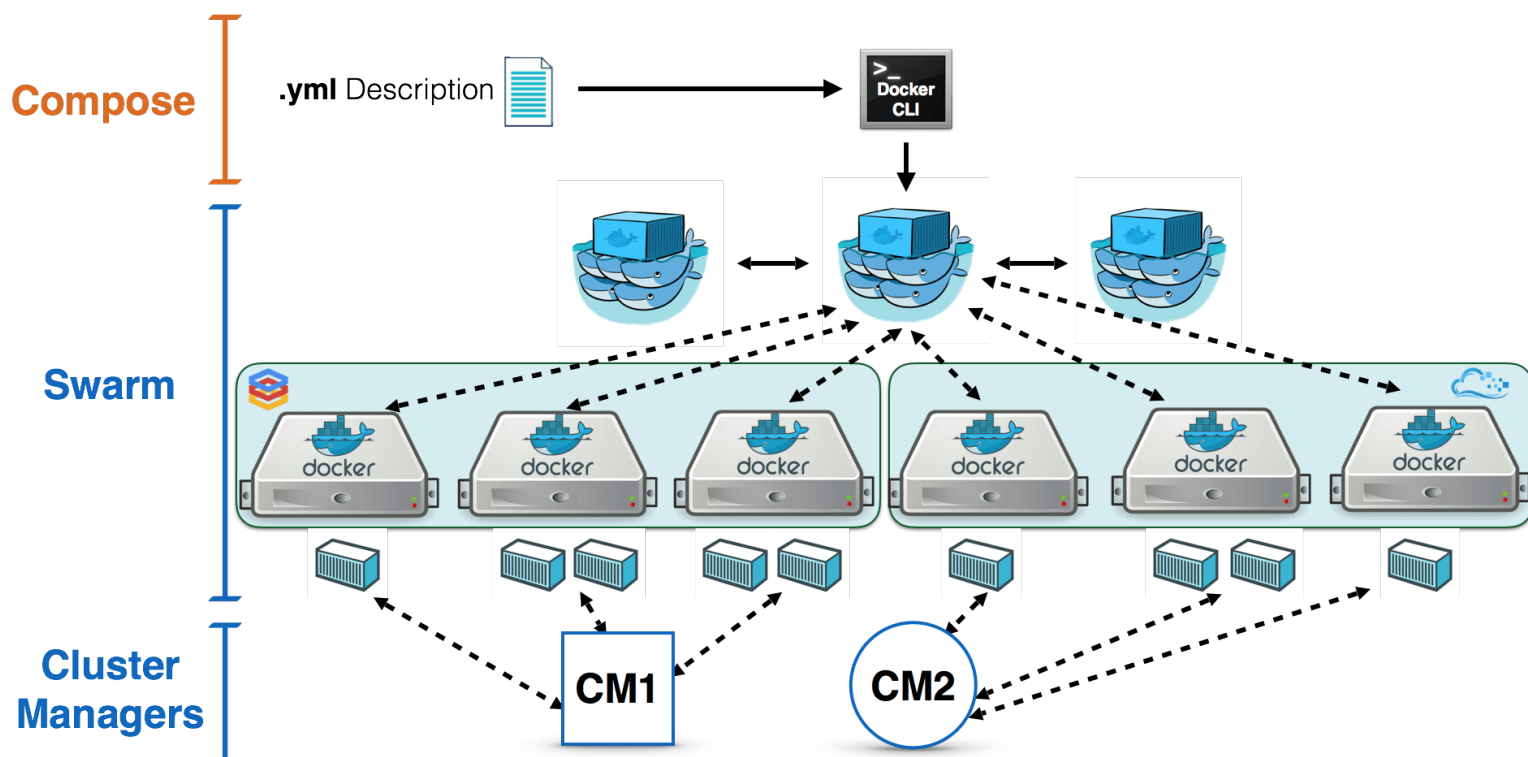
Reference

[Docker Swarm Cheatsheet-1](#)

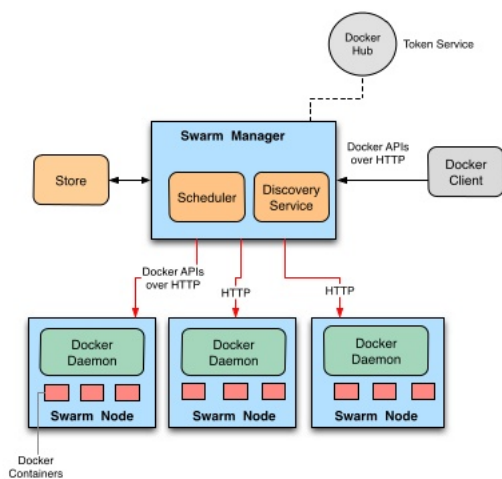
[Docker Swarm Cheatsheet-2](#)

[docker-machine](#) can be cloud or local VM. [Getting started with swarm mode](#)
[How to Create a Cluster of Docker Containers with Docker Swarm and DigitalOcean on Ubuntu 16.04](#)

Architecture



Docker Swarm Architecture - Exploded



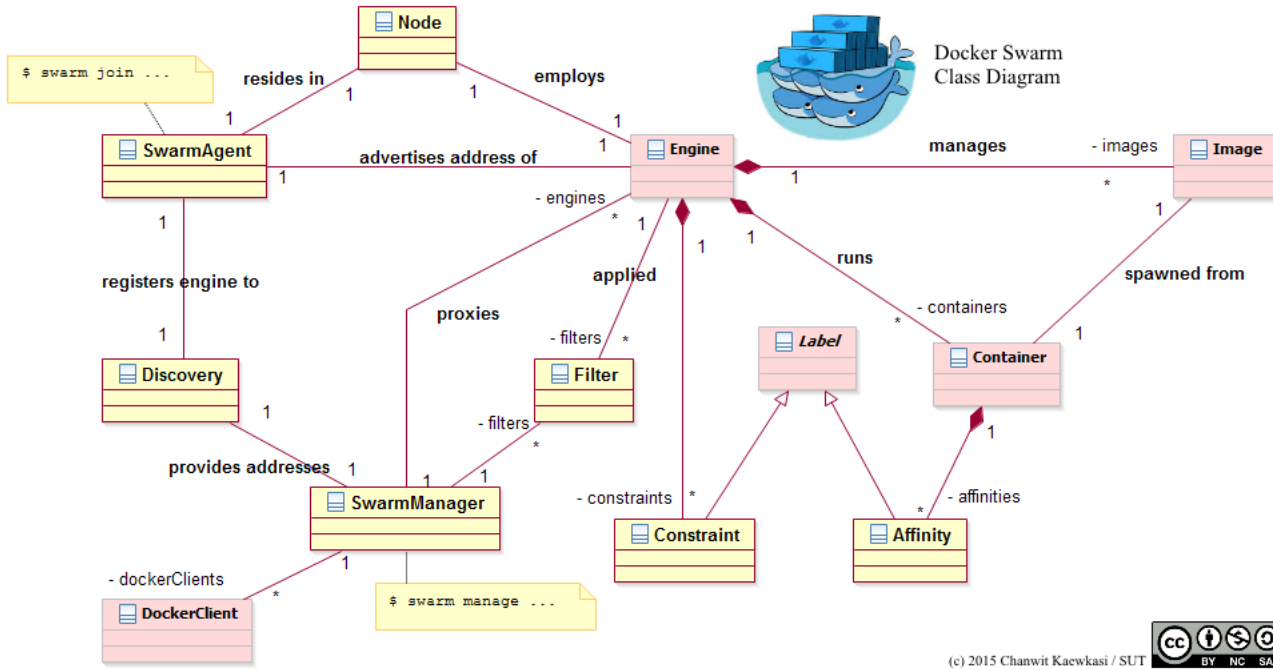
Features

highlights

- Cluster management
- Declarative service model: Docker Engine uses a declarative approach to let you define the desired state of the various services in your application stack.
- Scaling; task can be consider like POTs in Kubernetes.

- Maintain desired state via reconciliation.
- Multi-host networking: You can specify an overlay network for your services.
- Service discovery: Swarm manager nodes assign each service in the swarm a unique DNS name and load balances running containers.
- Load balancing: You can expose the ports for services to an external load balancer.
- Secure by default.
- Rolling updates the service.

UML



Initialize the Swarm Manager at Master Node

```

$ clear
HyprIoTOS/armv7: pirate@rpi-m1 in ~
$ docker swarm init
Swarm initialized: current node (g4vy8ewdl9ci8ct0lheqpxjbg) is now a manager.

To add a worker to this swarm, run the following command:

    docker swarm join \
    --token SWMTKN-1-0obaaqul2si12iu0mx1xwmoa0ig1frdsnjtyalpxnk7c1p434e-9l9sx9idajcloasfdw4a01g2y \
    192.168.1.88:2377

To add a manager to this swarm, run 'docker swarm join-token manager' and follow the instructions.

HyprIoTOS/armv7: pirate@rpi-m1 in ~

```

Add 3 Worker Nodes to Join Swarm

```

HyprIoTOS/armv7: pirate@rpi-s1 in ~
$ docker swarm join \
> --token SWMTKN-1-0obaaqul2si12iu0mx1xwmoa0ig1frdsnjtyalpxnk7c1p434e-9l9sx9idajcloasfdw4a01g2y \
> 192.168.1.88:2377
This node joined a swarm as a worker.
HyprIoTOS/armv7: pirate@rpi-s1 in ~
$

```

Review Swarm Info

```
HyprIoTOS/armv7: pirate@rpi-m1 in ~
$ docker info
Containers: 4
  Running: 0
  Paused: 0
  Stopped: 4
Images: 17
Server Version: 17.03.0-ce
Storage Driver: overlay2
  Backing Filesystem: extfs
  Supports d_type: true
  Native Overlay Diff: true
Logging Driver: json-file
Cgroup Driver: cgroupfs
Plugins:
  Volume: local
  Network: bridge host macvlan null overlay
Swarm: active <<-----
  NodeID: g4vy8ewdl9ci8ct0lheqpxjbg
  Is Manager: true
  ClusterID: pyrvg5qzok221w7pfzetgcox1
  Managers: 1
  Nodes: 4
  Orchestration:
    Task History Retention Limit: 5
  Raft:
    Snapshot Interval: 10000
    Number of Old Snapshots to Retain: 0
    Heartbeat Tick: 1
    Election Tick: 3
  Dispatcher:
    Heartbeat Period: 5 seconds
  CA Configuration:
    Expiry Duration: 3 months
  Node Address: 192.168.1.88
  Manager Addresses:
    192.168.1.88:2377
Runtimes: runc
Default Runtime: runc
Init Binary: docker-init
containerd version: 977c511eda0925a723debd94d09459af49d082a
runc version: a01dafd48bc1c7cc12bdb01206f9fea7dd6feb70
init version: 949e6fa
Kernel Version: 4.4.50-hyprIoTOS-v7+
Operating System: Raspbian GNU/Linux 8 (jessie)
OSType: linux
Architecture: armv7l
CPUs: 4
Total Memory: 861.9 MiB
Name: rpi-m1
ID: ZGHJ:GLKM:QIHZ:QKIZ:ZQAG:T7VW:STPB:3JX7:ESHQ:Y2ZQ:UDPA:NCRD
Docker Root Dir: /var/lib/docker
Debug Mode (client): false
Debug Mode (server): false
Registry: https://index.docker.io/v1/
Experimental: false
Insecure Registries:
  127.0.0.0/8
Live Restore Enabled: false
HyprIoTOS/armv7: pirate@rpi-m1 in ~
$
```

List of Nodes

```

HypriotOS/armv7: pirate@rpi-m1 in ~
$ docker node ls
ID                                HOSTNAME  STATUS  AVAILABILITY  MANAGER STATUS
g4vy8ewdl9ci8ct0lheqpxjbg *    rpi-m1   Ready   Active        Leader
u5oktv61epxq0jui2oq1iaytj      rpi-s2   Ready   Active
y7umdundbo2zm5cljz1wd708i       rpi-s3   Ready   Active
ywf90mm103qd28r6a0upldtu1       rpi-s1   Ready   Active
HypriotOS/armv7: pirate@rpi-m1 in ~

```

Running Services in the Docker Swarm

No service yet.

```

HypriotOS/armv7: pirate@rpi-m1 in ~
$ docker node ps
ID  NAME  IMAGE  NODE  DESIRED STATE  CURRENT STATE  ERROR  PORTS
HypriotOS/armv7: pirate@rpi-m1 in ~
$ docker service ls
ID  NAME  MODE  REPLICAS  IMAGE
HypriotOS/armv7: pirate@rpi-m1 in ~
$

```

Deploy it.

```

HypriotOS/armv7: pirate@rpi-m1 in ~
$ docker service create -p 80:80 --name webserver nginx
0eyzq92p6oduhcu56arphi23o
HypriotOS/armv7: pirate@rpi-m1 in ~
$ docker node ps
ID                                NAME      IMAGE           NODE  DESIRED STATE  CURRENT STATE           ERROR  PORTS
vy1gysrar89  webserver.1  nginx:latest    rpi-m1  Running        Preparing 26 seconds ago
HypriotOS/armv7: pirate@rpi-m1 in ~
$ docker service ls
ID  NAME      MODE      REPLICAS  IMAGE
0eyzq92p6odu  webserver  replicated  0/1        nginx:latest
HypriotOS/armv7: pirate@rpi-m1 in ~

```

```

HypriotOS/armv7: pirate@rpi-m1 in ~
$ docker service ls
ID  NAME  MODE  REPLICAS  IMAGE
HypriotOS/armv7: pirate@rpi-m1 in ~
$ docker network ls
NETWORK ID          NAME                DRIVER            SCOPE
647353628c83        bridge              bridge             local
698f31954f63        docker_gwbridge     bridge             local
ce2a385643c4        host                host               local
ifah9kmkvht2        ingress             overlay            swarm
a2800c300829        none                null               local
HypriotOS/armv7: pirate@rpi-m1 in ~
$ docker network create \
> --driver overlay \
> --subnet 10.10.1.0/24 \
> --opt encrypted \
> services
e3npyw8gejcyr5dkijqoyqwkp
HypriotOS/armv7: pirate@rpi-m1 in ~
$ docker network ls
NETWORK ID          NAME                DRIVER            SCOPE
647353628c83        bridge              bridge             local
698f31954f63        docker_gwbridge     bridge             local
ce2a385643c4        host                host               local
ifah9kmkvht2        ingress             overlay            swarm
a2800c300829        none                null               local
e3npyw8gejcy        services            overlay            swarm
HypriotOS/armv7: pirate@rpi-m1 in ~

```

why the service is not running?? `` HypriotOS/armv7: pirate@rpi-m1 in ~ \$ docker service create \


```
--replicas 2 \ --name nginx \ --network services \ --publish 80:80 \ nginx y50id9yg99z3ehada3mhgga1b HypriotOS/armv7: pirate@rpi-m1 in ~ $ docker service ls ID NAME MODE REPLICAS IMAGE y50id9yg99z3 nginx replicated 0/2 nginx:latest HypriotOS/armv7: pirate@rpi-m1 in ~ ``
```

do it again.

```
HypriotOS/armv7: pirate@rpi-m1 in ~ $ docker node ls ID HOSTNAME STATUS AVAILABILITY MANAGER STATUS g4vy8ewdl9ci8ct0lheqpxjbg * rpi-m1 Ready Service is running at rpi-s2.
```

```
HypriotOS/armv7: pirate@rpi-m1 in ~
$ docker service ls
ID            NAME    MODE     REPLICAS  IMAGE
r25uzlw259zs  ping   replicated 1/1       hypriot/rpi-alpine-scratch:latest
ttsu5t4sm164  web    replicated 0/2       nginx:latest
HypriotOS/armv7: pirate@rpi-m1 in ~
$
```

```
HypriotOS/armv7: pirate@rpi-s2 in ~
$ docker ps
CONTAINER ID   IMAGE                                     COMMAND                  CREATED
4b22f39ebc1f   hypriot/rpi-alpine-scratch@sha256:708171e6a1bd7c60a0ec9a5657900ada854bf14623be053f5865f918e0e2691c  "ping 8.8.8.8"         12 minutes ago
HypriotOS/armv7: pirate@rpi-s2 in ~
```

It is no at rpi-s3 or rpi-s1. scale up to 3;

```
HypriotOS/armv7: pirate@rpi-m1 in ~
$ docker service update --replicas 3 ping
ping
HypriotOS/armv7: pirate@rpi-m1 in ~
$ docker service ls
ID            NAME    MODE     REPLICAS  IMAGE
r25uzlw259zs  ping   replicated 1/3       hypriot/rpi-alpine-scratch:latest
ttsu5t4sm164  web    replicated 0/2       nginx:latest
HypriotOS/armv7: pirate@rpi-m1 in ~
$
```

still

```
HypriotOS/armv7: pirate@rpi-s2 in ~
$ docker ps
CONTAINER ID   IMAGE                                     COMMAND                  CREATED
4b22f39ebc1f   hypriot/rpi-alpine-scratch@sha256:708171e6a1bd7c60a0ec9a5657900ada854bf14623be053f5865f918e0e2691c  "ping 8.8.8.8"         22 minutes ago
HypriotOS/armv7: pirate@rpi-s2 in ~
$
```

Monitoring

- [Visualize your Raspberry Pi containers with Portainer or UI for Docker](#)
- [Deploying an IoT Swarm with Docker Machine](#)

Physical IO

- [Let's get physical with Docker on the Raspberry Pi](#)
- [Wiring Pi - GPIO Interface library for the Raspberry Pi](#)

Summary

Cheat Sheet

Docker Cheat Sheet



Initialize swarm mode and listen on a specific interface
`docker swarm init --advertise-addr 10.1.0.2`

Join an existing swarm as a manager node
`docker swarm join --token <manager-token> 10.1.0.2:2377`

Join an existing swarm as a worker node
`docker swarm join --token <worker-token> 10.1.0.2:2377`

List the nodes participating in a swarm
`docker node ls`

Create a service from an image exposed on a specific port and deploy 3 instances
`docker service create --replicas 3 -p 80:80 --name web nginx`

List the services running in a swarm
`docker service ls`

Scale a service
`docker service scale web=5`

List the tasks of a service
`docker service tasks web`

RUN

`docker run`
--rm remove container automatically after it exits
-it connect the container to terminal
--name web name the container
-p 5000:80 expose port 5000 externally and map to port 80
-v ~/dev:/code create a host mapped volume inside the container
alpine:3.4 the image from which the container is instantiated
/bin/sh the command to run inside the container

Stop a running container through SIGTERM
`docker stop web`

Stop a running container through SIGKILL
`docker kill web`

Create an overlay network and specify a subnet
`docker network create --subnet 10.1.0.0/24 --gateway 10.1.0.1 -d overlay mynet`

List the networks
`docker network ls`

List the running containers
`docker ps`

Delete all running and stopped containers
`docker rm -f $(docker ps -aq)`

Create a new bash process inside the container and connect it to the terminal
`docker exec -it web bash`

Print the last 100 lines of a container's logs
`docker logs --tail 100 web`

BUILD

Build an image from the Dockerfile in the current directory and tag the image
`docker build -t myapp:1.0 .`

List all images that are locally stored with the Docker engine
`docker images`

Delete an image from the local image store
`docker rmi alpine:3.4`

SHIP

Pull an image from a registry
`docker pull alpine:3.4`

Retag a local image with a new image name and tag
`docker tag alpine:3.4 myrepo/myalpine:3.4`

Log in to a registry (the Docker Hub by default)
`docker login my.registry.com:8000`

Push an image to a registry
`docker push myrepo/myalpine:3.4`

CLI

swarm init swarm join service create service inspect service ls service rm service scale service ps service update